



FROM

MODERN ALGORITHMS WORKSHOP

Parallel Algorithms

Prof. Charles E. Leiserson

Dr. Tao B. Schardl

September 19, 2018

Outline

- Introduction
- Cilk Model
- Detecting Nondeterminism
- What Is Parallelism?
- Scheduling Theory Primer
- *Lunch Break*
- Analysis of Parallel Loops
- Case Study: Matrix Multiplication
- Case Study: Jaccard Similarity
- Post-Moore Software

WHAT IS PARALLELISM?



Execution Model

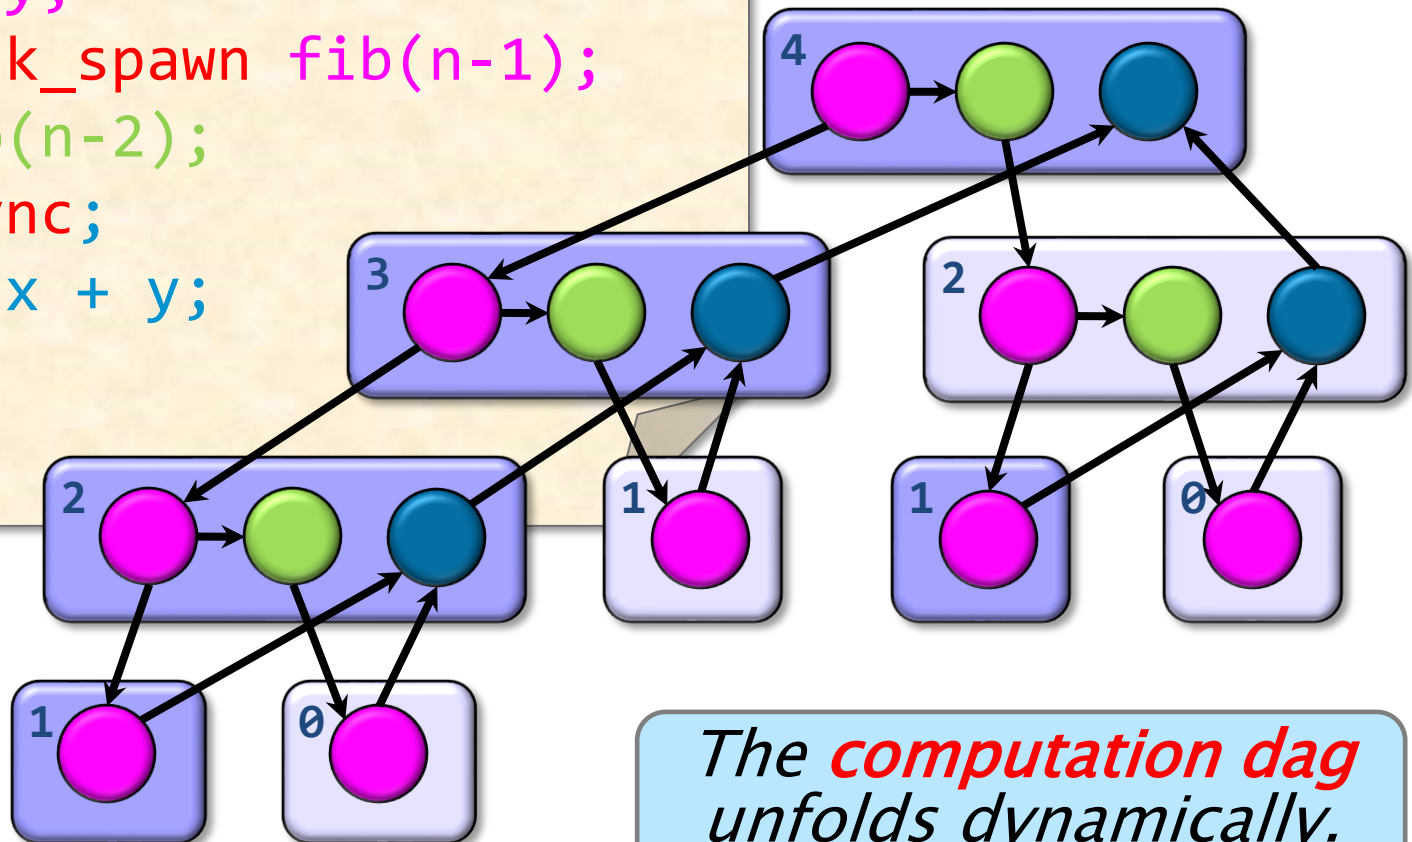
```
int fib (int n) {  
    if (n < 2) return n;  
    else {  
        int x, y;  
        x = cilk_spawn fib(n-1);  
        y = fib(n-2);  
        cilk_sync;  
        return x + y;  
    }  
}
```

Example:
fib(4)

Execution Model

```
int fib (int n) {  
  if (n < 2) return n;  
  else {  
    int x, y;  
    x = cilk_spawn fib(n-1);  
    y = fib(n-2);  
    cilk_sync;  
    return x + y;  
  }  
}
```

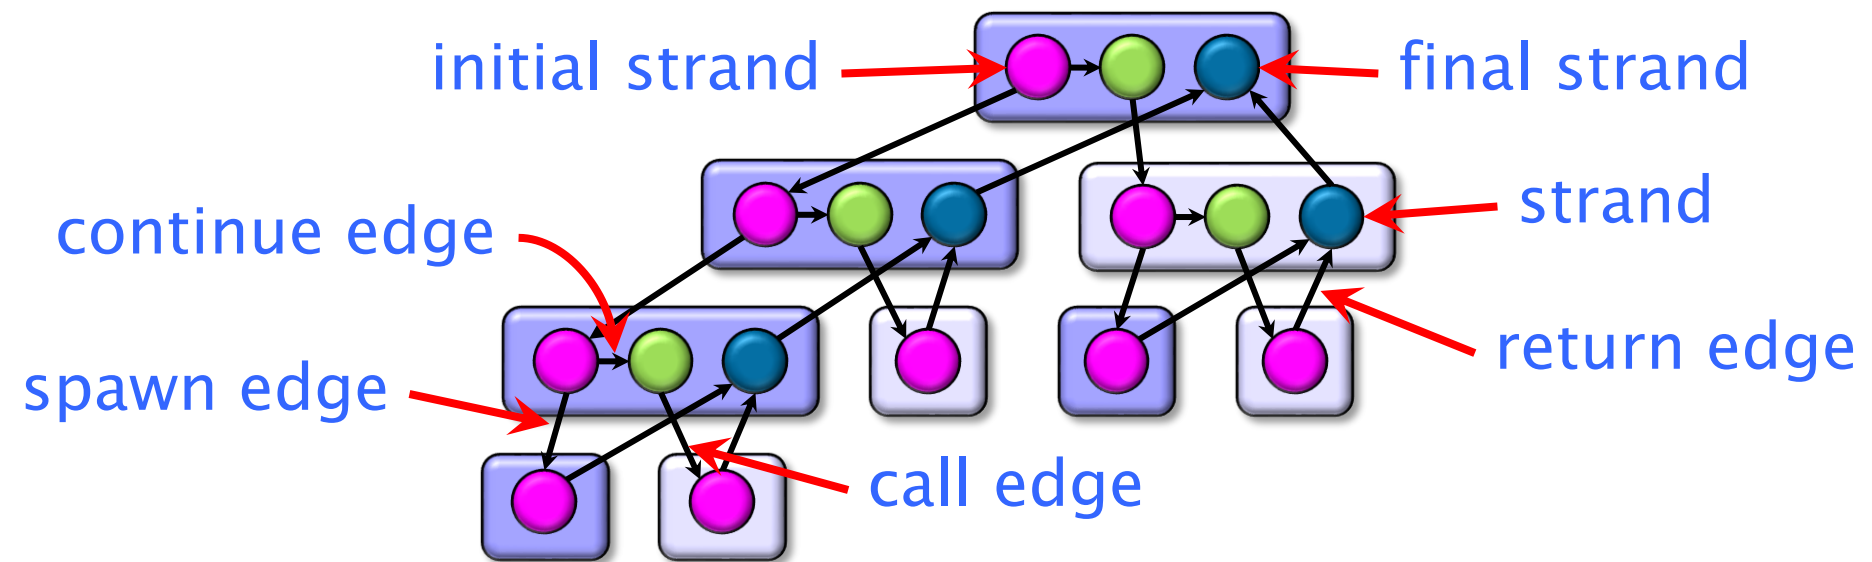
Example:
fib(4)



“Processor
oblivious”

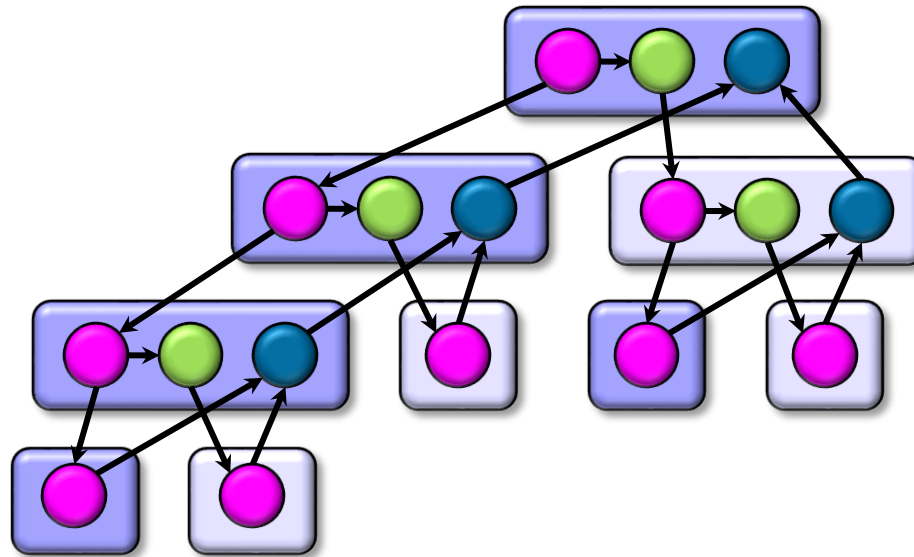
*The **computation dag**
unfolds dynamically.*

Computation Dag



- A parallel instruction stream is a dag $G = (V, E)$.
- Each vertex $v \in V$ is a **strand**: a sequence of instructions not containing a spawn, sync, or return from a spawn.
- An edge $e \in E$ is a **spawn**, **call**, **return**, or **continue** edge.
- Loop parallelism (**cilk_for**) is converted to spawns and syncs using recursive divide-and-conquer.

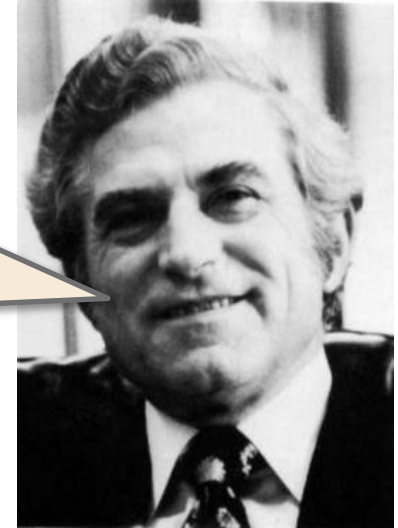
How Much Parallelism?



Assuming that each strand executes in unit time, what is the **parallelism** of this computation?

Amdahl's "Law"

If 50% of your application is parallel and 50% is serial, you can't get more than a factor of 2 speedup, no matter how many processors it runs on.

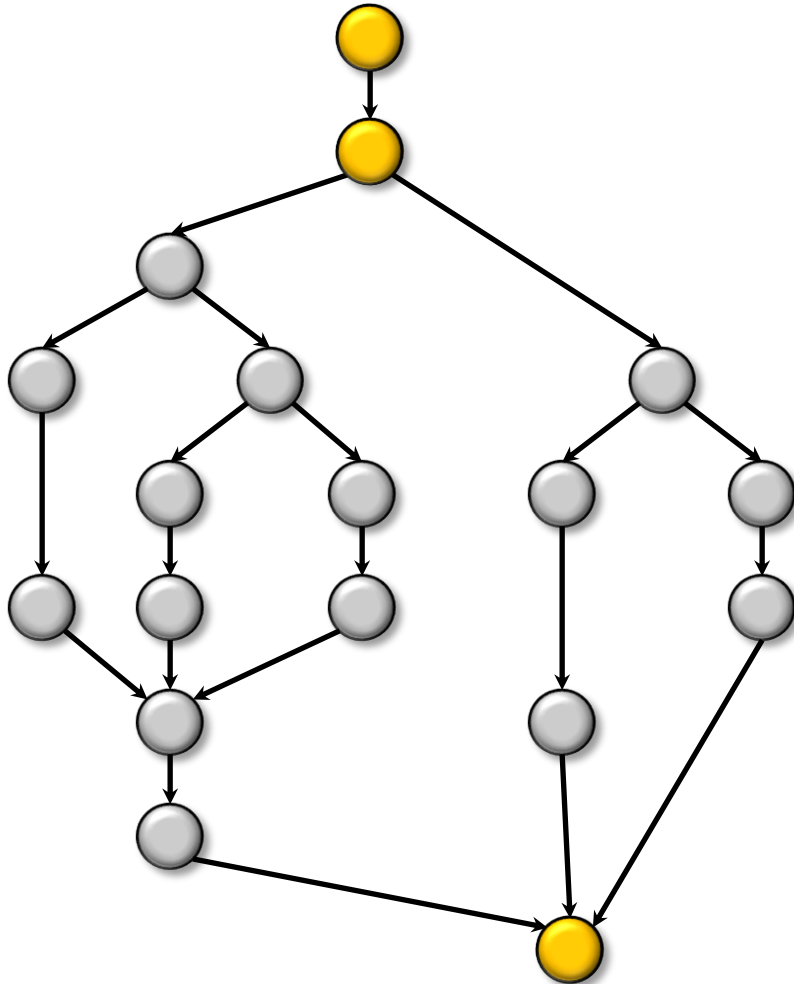


Gene M. Amdahl

In general, if a fraction α of an application must be run serially, the speedup can be at most $1/\alpha$.

Quantifying Parallelism

What is the **parallelism** of this computation?

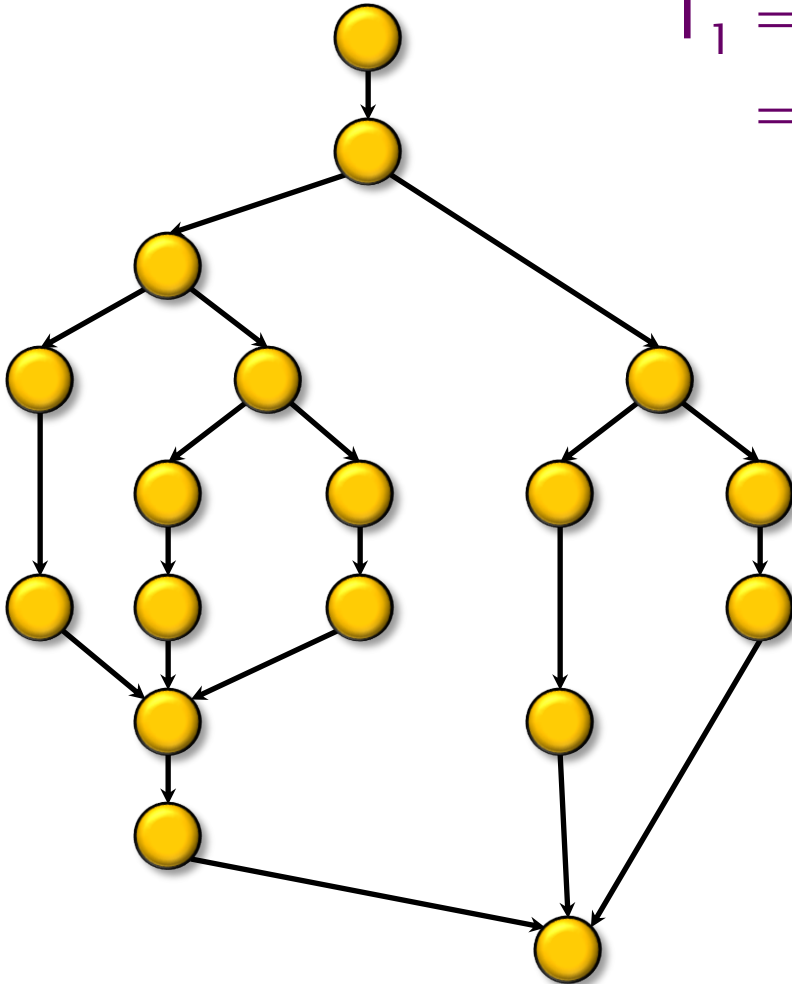


Amdahl's Law says that since the serial fraction is $3/18 = 1/6$, the speedup is upper-bounded by 6.

Performance Measures

T_p = execution time on P processors

T_1 = work
= 18

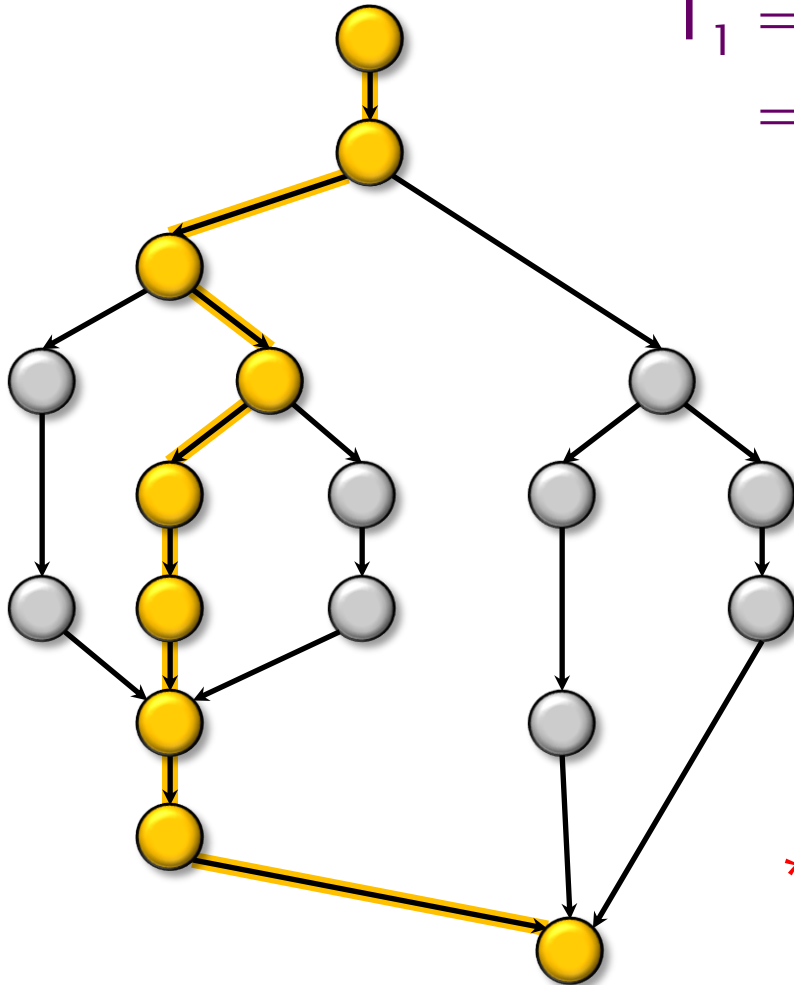


Performance Measures

T_p = execution time on P processors

$$T_1 = \text{work} = 18$$

$$T_\infty = \text{span}^* \\ = 9$$



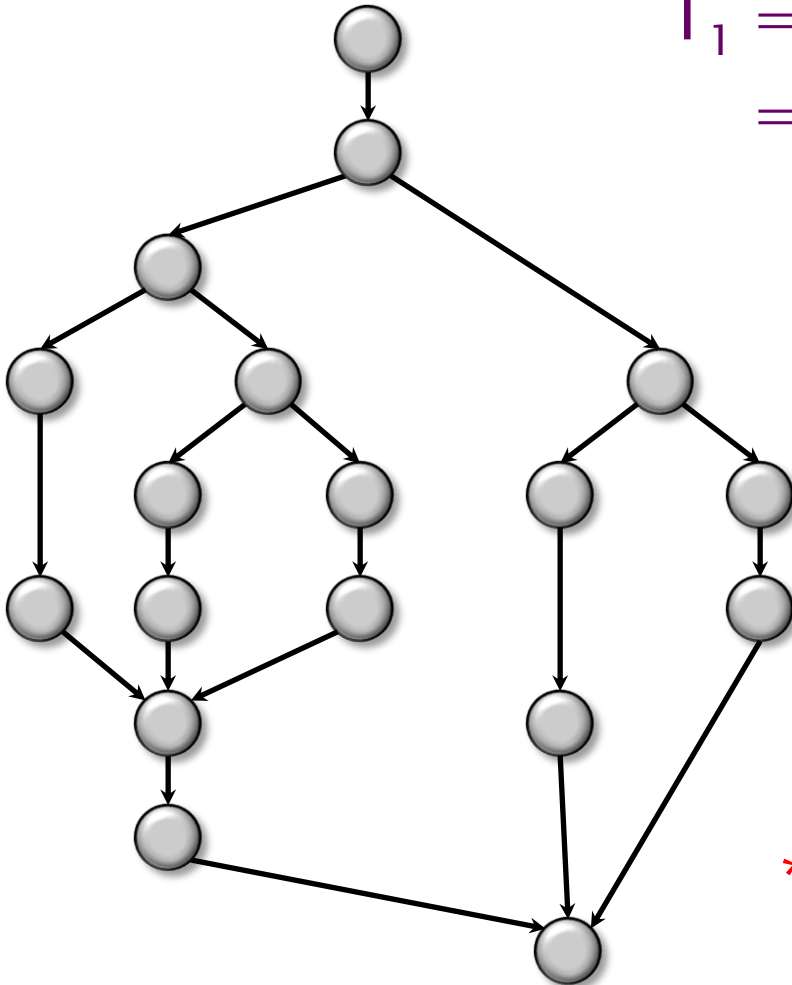
- * Also called **critical-path length** or **computational depth**.

Performance Measures

T_p = execution time on P processors

$$T_1 = \text{work} = 18$$

$$T_{\infty} = \text{span}^* \\ = 9$$



WORK LAW

- $T_p \geq T_1/P$

SPAN LAW

- $T_p \geq T_\infty$

- * Also called **critical-path length** or **computational depth**.

Speedup

Definition. $T_1/T_P = \text{speedup}$ on P processors.

-
- If $T_1/T_P < P$, we have **sublinear speedup**.
 - If $T_1/T_P = P$, we have **(perfect) linear speedup**.
 - If $T_1/T_P > P$, we have **superlinear speedup**, which is not possible in this simple performance model, because of the **WORK LAW** $T_P \geq T_1/P$.
-

Parallelism

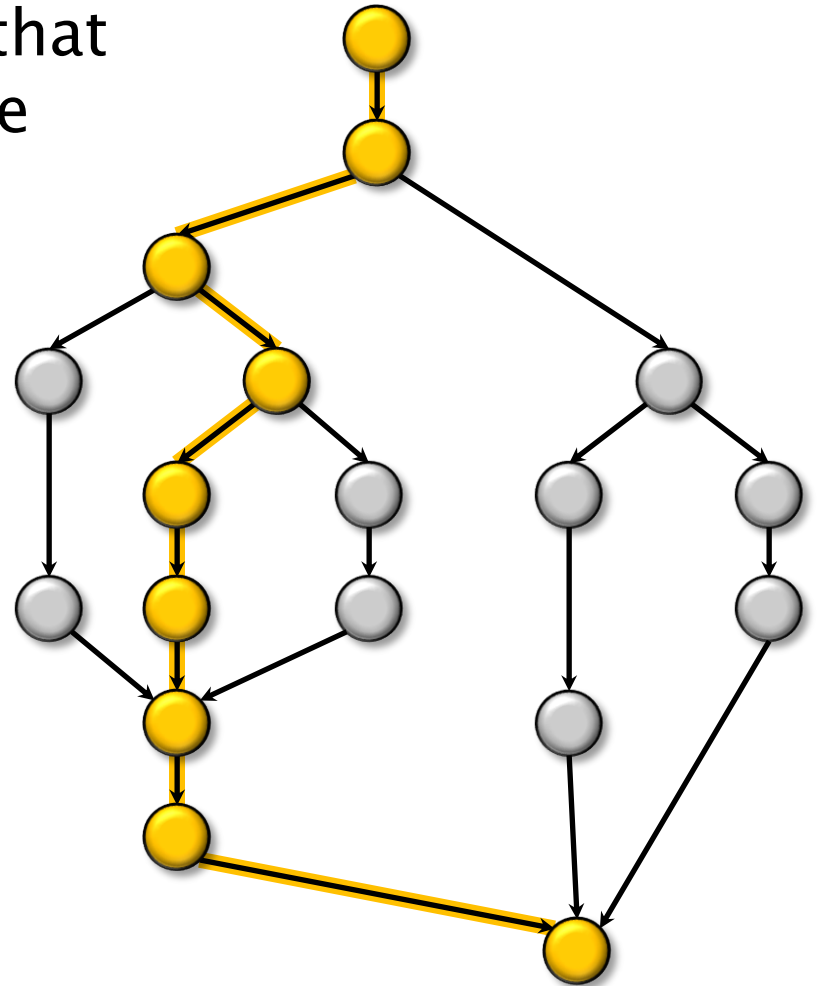
Because the **SPAN LAW** dictates that $T_p \geq T_\infty$, the maximum possible speedup given T_1 and T_∞ is

$T_1/T_\infty = \text{parallelism}$

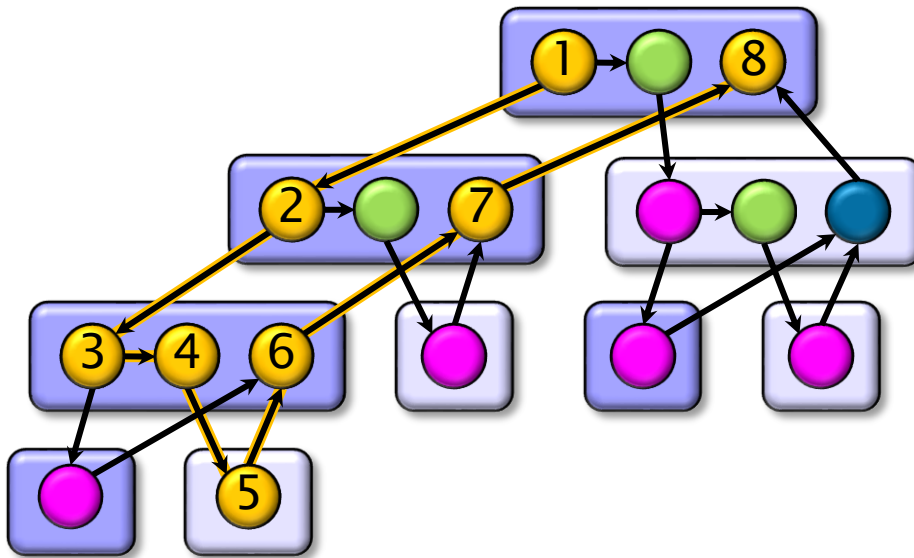
= the average
amount of work
per step along
the span

= $18/9$

= 2 .



Example: fib(4)



Assume for simplicity that each strand in **fib(4)** takes unit time to execute.

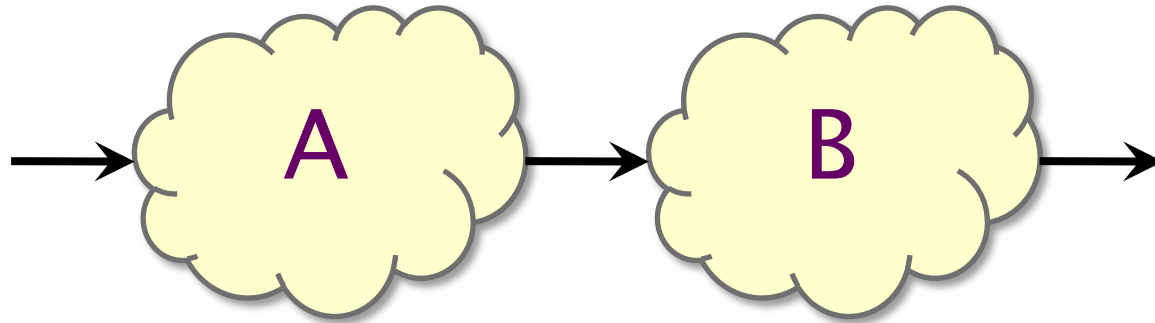
Work: $T_1 = 17$

Span: $T_\infty = 8$

Parallelism: $T_1/T_\infty = 2.125$

Using more than 2 processors guarantees that some processors will be idle.

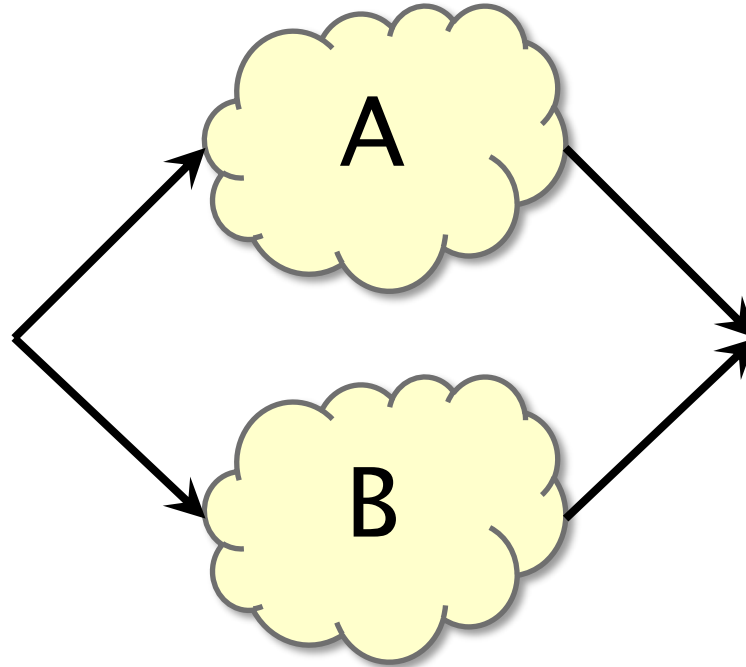
Quiz: Series Composition



Work: $T_1(A \cup B) = T_1(A) + T_1(B)$

Span: $T_\infty(A \cup B) = T_\infty(A) + T_\infty(B)$

Quiz: Parallel Composition



Work: $T_1(A \cup B) = T_1(A) + T_1(B)$

Span: $T_\infty(A \cup B) = \max\{T_\infty(A), T_\infty(B)\}$

HANDS-ON: THE CILKSCALE SCALABILITY ANALYZER



Quicksort Analysis

Example: Parallel quicksort

```
void quicksort(int64_t *left, int64_t *right)
{
    int64_t *middle;
    if (left == right) return;
    middle = partition(left, right);
    cilk_spawn quicksort(left, middle);
    quicksort(middle + 1, right);
    cilk_sync;
}
```

Analyze the sorting of 1,000,000 numbers.

★ ★ ★ *Guess the parallelism!* ★ ★ ★

Cilkscale Scalability Analyzer

- The Tapir/LLVM compiler provides a **scalability analyzer** called **Cilkscale**.
- Like the Cilksan race detector, Cilkscale uses **compiler-instrumentation** to analyze a serial execution of a program.
- Cilkscale computes **work** and **span** to derive upper bounds on parallel performance.

Run Cilkscale on QSort

1. Use make to compile `qsort`:

```
$ cd qsort  
$ make
```

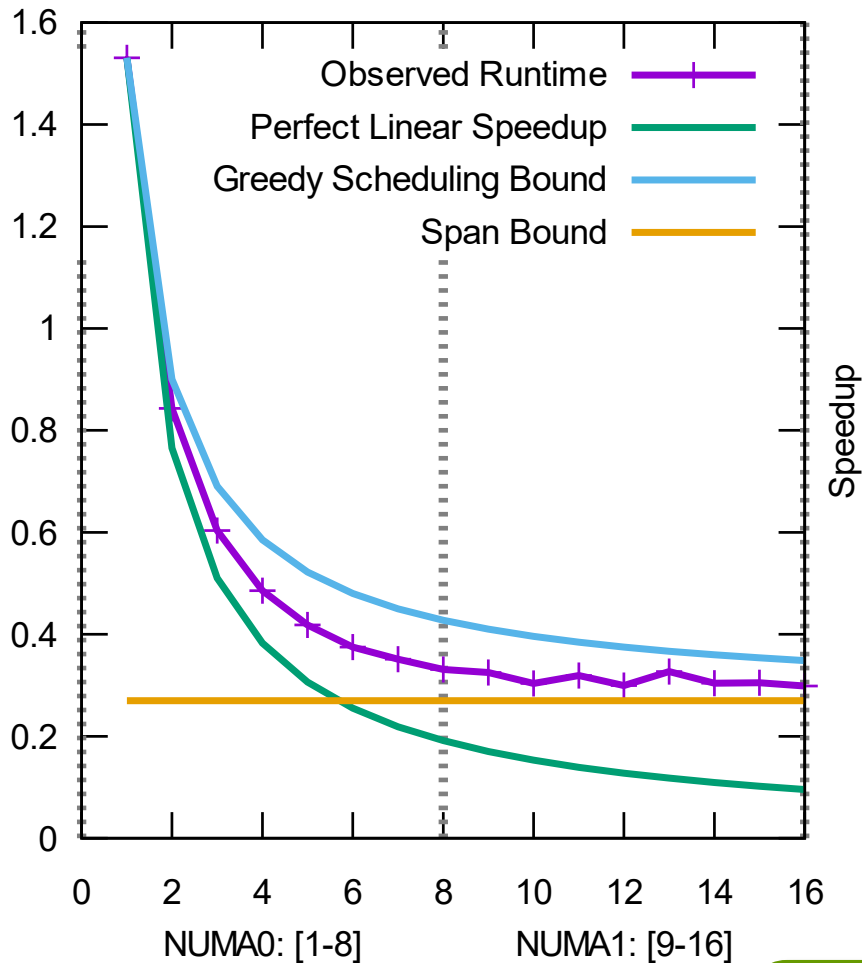
2. Run the Cilkscale script on `qsort` of 1,000,000 elements:

```
$ cilkscale ./qsort 1000000
```

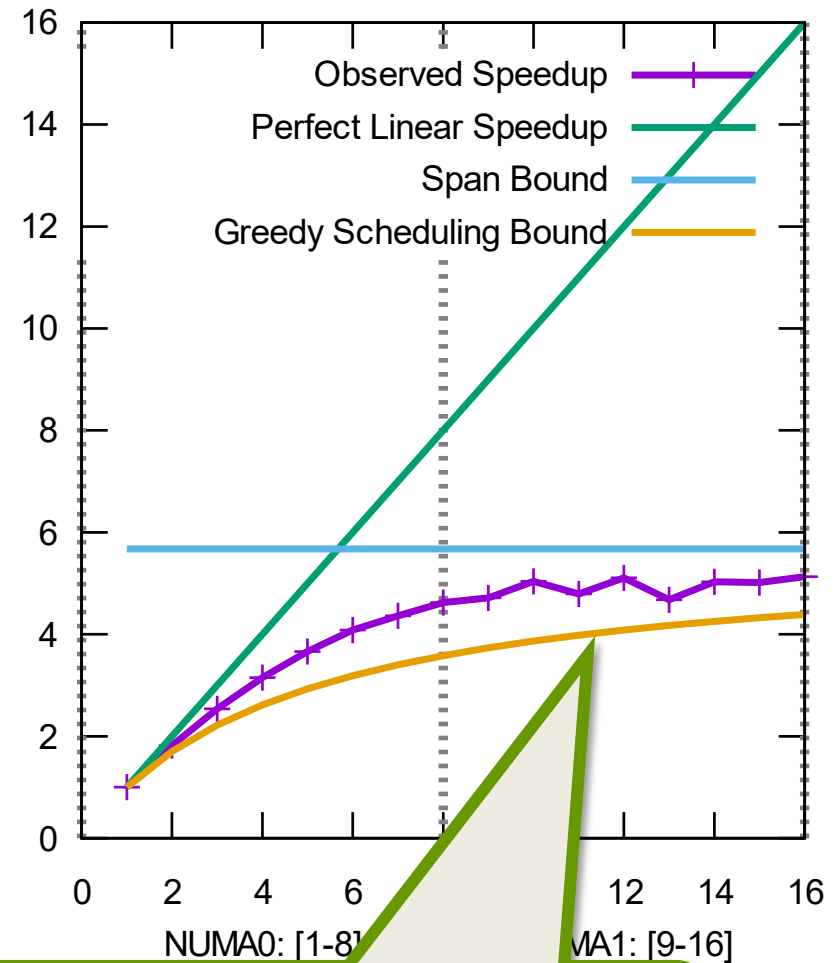
3. Browse the files in your Jupyter notebook to find the output file `qsort-1000000.svg`, and open that file in your browser:
 - Click the checkbox next to the file.
 - Click the “View” button at the top of the window.

Cilkscale Output

Execution Time: ./qsort 1000000



Speedup: ./qsort 1000000



We will derive this bound in the next section.

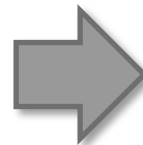
Theoretical Analysis

Example: Parallel quicksort

```
void quicksort(int64_t *left, int64_t *right)
{
    int64_t *middle;
    if (left == right) return;
    middle = partition(left, right);
    cilk_spawn quicksort(left, middle);
    quicksort(middle + 1, right);
    cilk_sync;
}
```

Expected work = $O(n \lg n)$

Expected span = $\Omega(n)$



Parallelism = $O(\lg n)$

Interesting Practical* Algorithms

Algorithm	Work	Span	Parallelism
Merge sort	$\Theta(n \lg n)$	$\Theta(\lg^3 n)$	$\Theta(n / \lg^2 n)$
Matrix multiplication	$\Theta(n^3)$	$\Theta(\lg n)$	$\Theta(n^3 / \lg n)$
Strassen	$\Theta(n^{\lg 7})$	$\Theta(\lg^2 n)$	$\Theta(n^{\lg 7} / \lg^2 n)$
LU-decomposition	$\Theta(n^3)$	$\Theta(n \lg n)$	$\Theta(n^2 / \lg n)$
Tableau construction	$\Theta(n^2)$	$\Theta(n^{\lg 3})$	$\Theta(n^{2-\lg 3})$
FFT	$\Theta(n \lg n)$	$\Theta(\lg^2 n)$	$\Theta(n / \lg n)$
Breadth-first search	$\Theta(E)$	$\Theta(D \lg V)$	$\Theta(E / D \lg V)$

*Cilk on 1 processor competitive with the best C.