



FROM

MODERN ALGORITHMS WORKSHOP

# Parallel Algorithms

Prof. Charles E. Leiserson

Dr. Tao B. Schardl

*September 19, 2018*

# Outline


- Introduction
- Cilk Model
- Detecting Nondeterminism
- What Is Parallelism?
- Scheduling Theory Primer
- *Lunch Break*
- Analysis of Parallel Loops
- Case Study: Matrix Multiplication
- Case Study: Jaccard Similarity
- Post-Moore Software

# SCHEDULING THEORY PRIMER



# Scheduling


- Cilk allows the programmer to express **potential parallelism** in an application.
- The Cilk **scheduler** maps strands onto processors dynamically at runtime.
- Since the theory of **distributed** schedulers is complicated, we'll explore the ideas with a **centralized** scheduler.



# Greedy Scheduling

**IDEA:** Do as much as possible on every step.

**Definition.** A strand is **ready** if all its predecessors have executed.




# Greedy Scheduling

**IDEA:** Do as much as possible on every step.

**Definition.** A strand is **ready** if all its predecessors have executed.

## Complete step

- $\geq P$  strands ready.
- Run any  $P$ .



# Greedy Scheduling

**IDEA:** Do as much as possible on every step.


**Definition.** A strand is **ready** if all its predecessors have executed.

## Complete step

- $\geq P$  strands ready.
- Run any  $P$ .

## Incomplete step

- $< P$  strands ready.
- Run all of them.




# Analysis of Greedy

**Theorem** [G68, B75, EZL89]. Any greedy scheduler achieves

$$T_p \leq T_1/P + T_\infty.$$

*Proof.*

- # complete steps  $\leq T_1/P$ , since each complete step performs  $P$  work.
- # incomplete steps  $\leq T_\infty$ , since each incomplete step reduces the span of the unexecuted dag by 1. ■





# Optimality of Greedy

**Corollary.** Any greedy scheduler achieves within a factor of 2 of optimal.

*Proof.* Let  $T_p^*$  be the execution time produced by the optimal scheduler. Since  $T_p^* \geq \max\{T_1/P, T_\infty\}$  by the WORK and SPAN LAWS, we have

$$\begin{aligned} T_p &\leq T_1/P + T_\infty \\ &\leq 2 \cdot \max\{T_1/P, T_\infty\} \\ &\leq 2T_p^* . \quad \blacksquare \end{aligned}$$

# Linear Speedup

**Corollary.** Any greedy scheduler achieves near-perfect linear speedup whenever  $T_1/T_\infty \gg P$ .

*Proof.* Since  $T_1/T_\infty \gg P$  is equivalent to  $T_\infty \ll T_1/P$ , the Greedy Scheduling Theorem gives us

$$\begin{aligned} T_P &\leq T_1/P + T_\infty \\ &\approx T_1/P. \end{aligned}$$

Thus, the speedup is  $T_1/T_P \approx P$ . ■

**Definition.** The quantity  $T_1/PT_\infty$  is called the parallel slackness.

# Cilk Performance

- Cilk's work-stealing scheduler achieves
  - $T_p = T_1/P + O(T_\infty)$  expected time (provably);
  - $T_p \approx T_1/P + T_\infty$  time (empirically).
- Near-perfect **linear speedup** as long as  $P \ll T_1/T_\infty$ .
- Instrumentation in Cilkscale allows you to measure  $T_1$  and  $T_\infty$ .

# Quiz on Scheduling Theory

Your team is performance-engineering a parallel program.\* Although the program will be run on a 512-processor ideal machine, you are testing the program on a similar but smaller 32-processor machine. One of the engineers proposes a change that improves the 1- and 32-processor performance of the code as follows:

	$T_1$	$T_{32}$
<i>Before</i>	2048	65
<i>After</i>	1024	40

Assume that the running times obey the equation  $T_p = T_1/P + T_\infty$ . Is the optimization wise? Justify.

---

\*This scenario really happened during the development of the ★Socrates chess-playing program, although the timing numbers have been simplified.

# Quiz Solution

	$T_1$	$T_{32}$
<i>Before</i>	2048	65
<i>After</i>	1024	40

*Before:*

$$65 = 2048/32 + T_\infty$$

$$\Rightarrow T_\infty = 1$$

$$\begin{aligned}\Rightarrow T_{512} &= 2048/512 + 1 \\ &= 5\end{aligned}$$

*After:*

$$40 = 1024/32 + T_\infty$$

$$\Rightarrow T_\infty = 8$$

$$\begin{aligned}\Rightarrow T_{512} &= 1024/512 + 8 \\ &= 10\end{aligned}$$

*Reject the optimization!*

# Quiz 2 on Scheduling Theory

Ben Bitdiddle measures the running time of his deterministic parallel program scheduled using a greedy scheduler on an ideal parallel computer with 4, 10, and 64 processors. Ben obtains the following running times:

Processors	Time
$T_4$	80
$T_{10}$	42
$T_{64}$	9

Argue that Ben messed up at least one of his measurements.

# Quiz 2 Solution

Processors	Time
$T_4$	80
$T_{10}$	42
$T_{64}$	9

*Work Law:*  $T_1 \leq 4 * T_4$   
 $= 320$

*Span Law:*  $T_\infty \leq T_{64}$   
 $= 9$

*Greedy scheduling on 10 processors:*  $T_{10} \leq 320/10 + 9$   
 $= 41$

But Ben measured  $T_{10} = 42!$  ■