



MIT COMPUTER SCIENCE AND ARTIFICIAL INTELLIGENCE LABORATORY



FROM

MODERN ALGORITHMS WORKSHOP

Parallel Algorithms

Prof. Charles E. Leiserson

Dr. Tao B. Schardl

*September 19, 2018*

# Outline

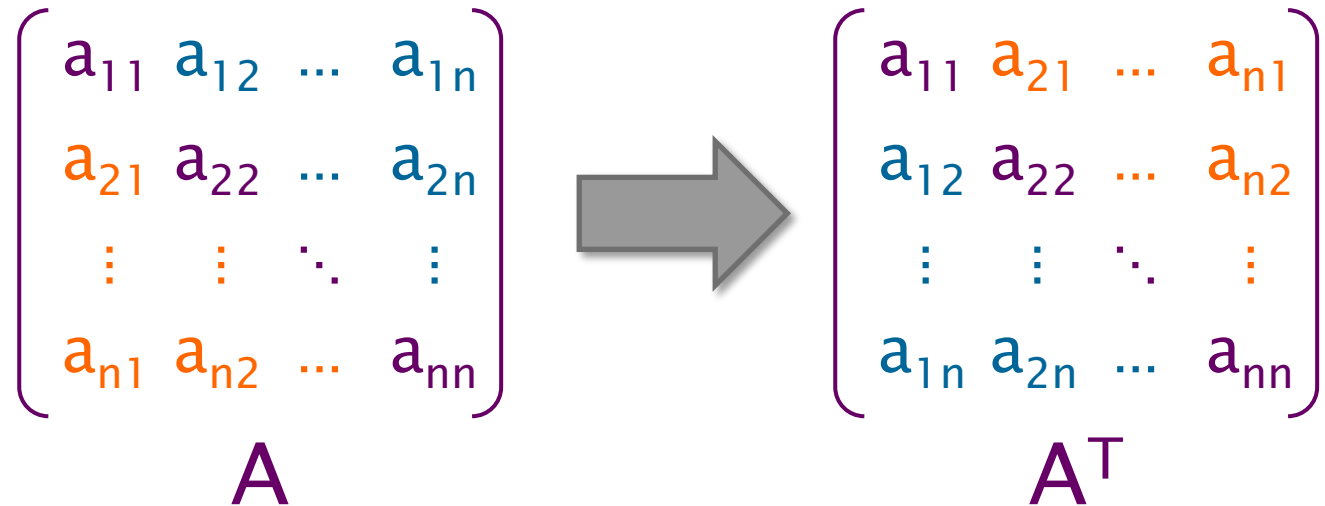
- Introduction
- Cilk Model
- Detecting Nondeterminism
- What Is Parallelism?
- Scheduling Theory Primer
- *Lunch Break*
- Analysis of Parallel Loops
- Case Study: Matrix Multiplication
- Case Study: Jaccard Similarity
- Post-Moore Software

# ANALYSIS OF PARALLEL LOOPS



# Loop Parallelism in Cilk

**Example:**  
In-place  
matrix  
transpose



The iterations of a  
**cilk\_for** loop  
execute in parallel.

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
```

# Implementation of Parallel Loops

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
```

## *Divide-and-conquer implementation*

The Tapir/LLVM compiler implements `cilk_for` loops this way at optimization level `-O1` or higher.

```
void recur(int lo, int hi) //half open
{
    if (hi > lo + 1) {
        int mid = lo + (hi - lo)/2;
        cilk_spawn recur(lo, mid);
                    recur(mid, hi);
        cilk_sync;
        return;
    }
    int i = lo;
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
...
recur(1, n);
```

# Implementation of Parallel Loops

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
```

**cilk\_for**  
loop control

*Divide-and-conquer  
implementation*

*lifted*  
loop body

```
void recur(int lo, int hi) //half open
{
    if (hi > lo + 1) {
        int mid = lo + (hi - lo)/2;
        cilk_spawn recur(lo, mid);
                    recur(mid, hi);
        cilk_sync;
        return;
    }
    int i = lo;
    for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
    ...
    recur(1, n);
}
```

# Execution of Parallel Loops

```
void recur(int lo, int hi) //half open  
{
```

```
    if (hi > lo + 1) {  
        int mid = lo + (hi - lo)/2;  
        cilk_spawn recur(lo, mid);  
                   recur(mid, hi);  
        cilk_sync;  
        return;  
    }
```

```
    int i = lo;  
    for (int j=0; j<i; ++j) {  
        double temp = A[i][j];  
        A[i][j] = A[j][i];  
        A[j][i] = temp;  
    }
```

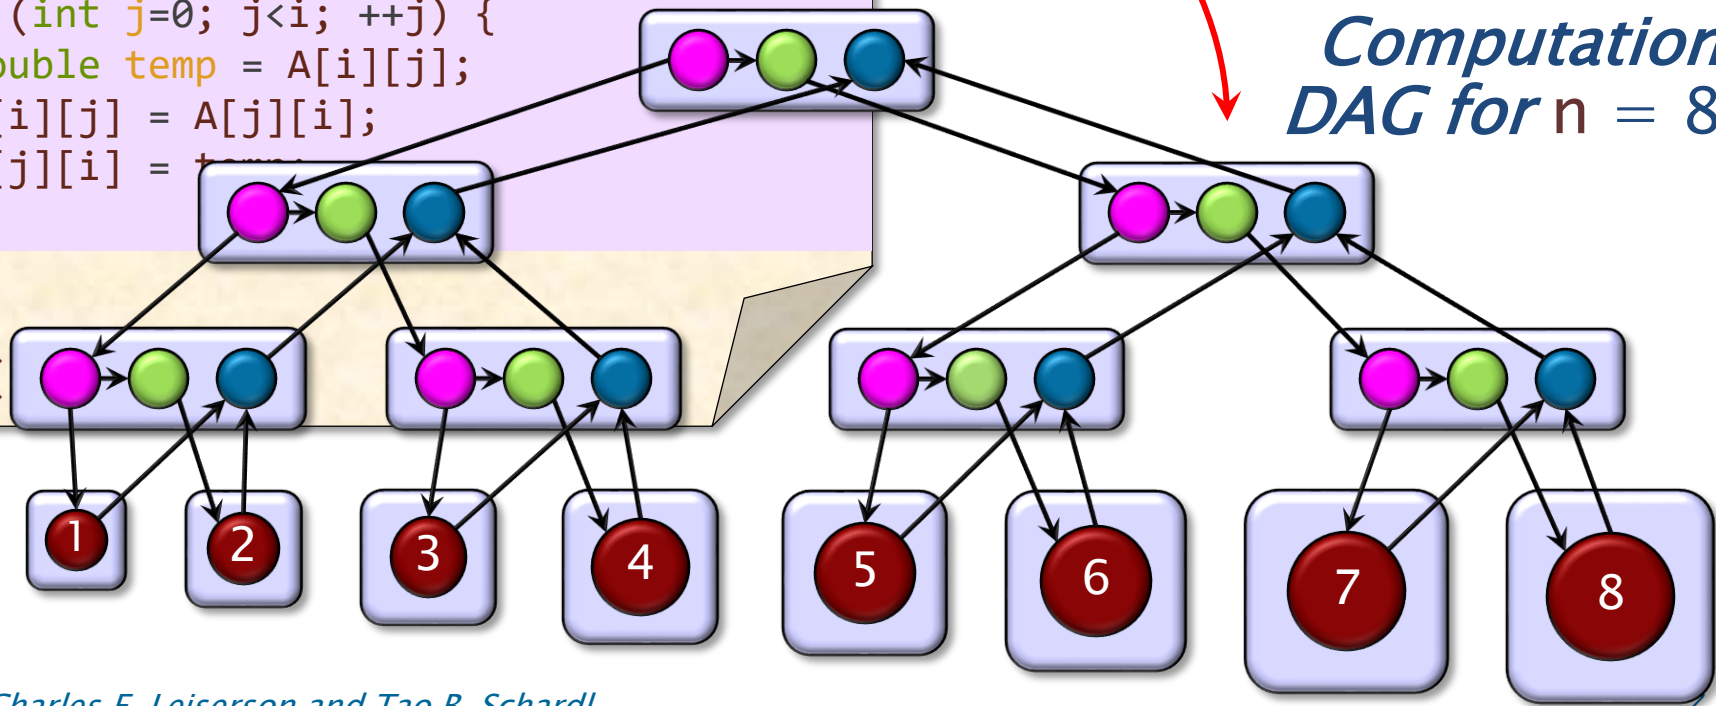
```
}
```

```
recur(
```

*Divide-and-conquer  
implementation*

*cilk\_for*  
loop control

*Computation  
DAG for n = 8*

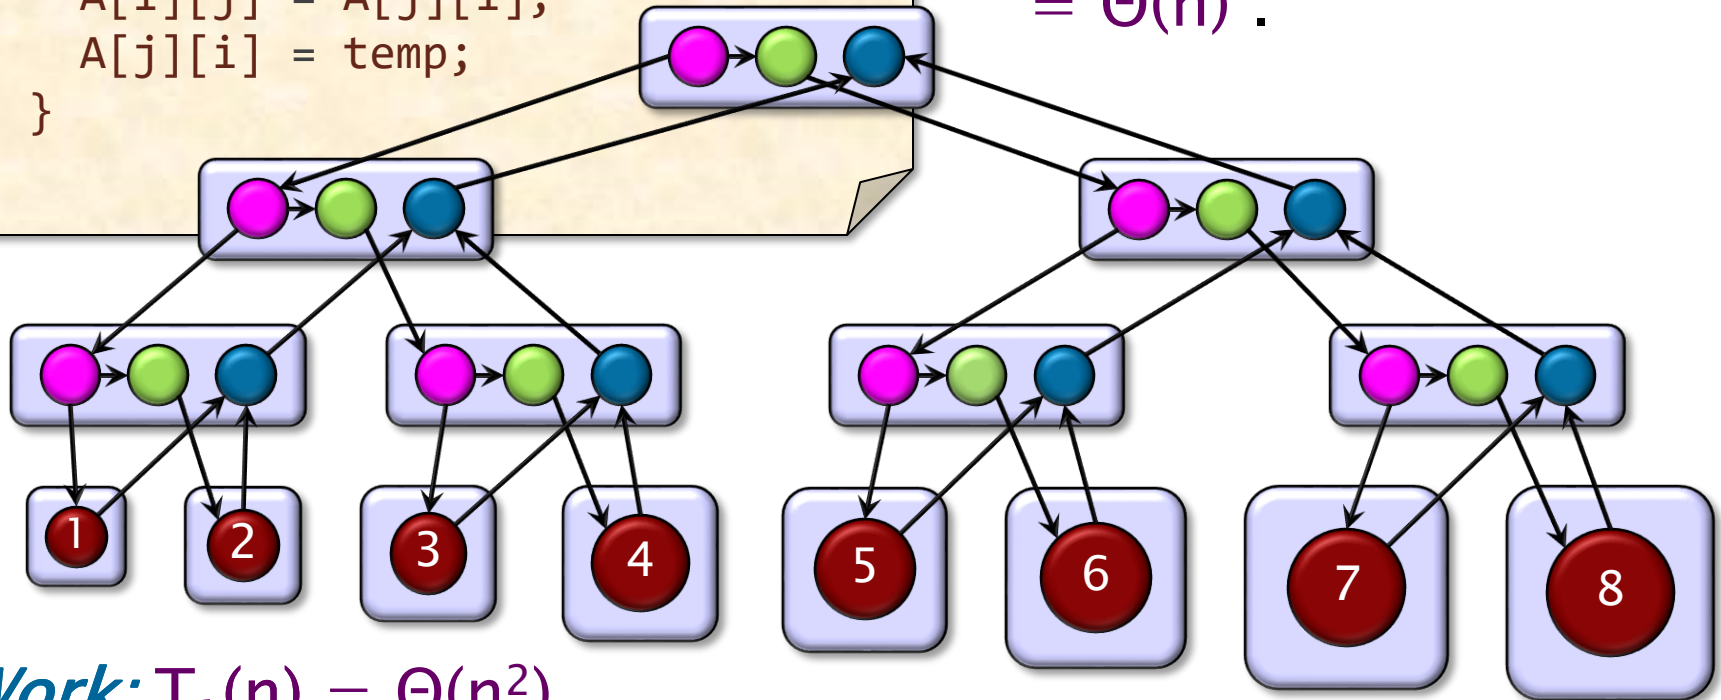


# Analysis of Parallel Loops

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
  for (int j=0; j<i; ++j) {
    double temp = A[i][j];
    A[i][j] = A[j][i];
    A[j][i] = temp;
  }
}
```

Span of loop control  
 $= \Theta(\lg n)$ .

Max span of body  
 $= \Theta(n)$ .



**Work:**  $T_1(n) = \Theta(n^2)$

**Span:**  $T_\infty(n) = \Theta(n + \lg n) = \Theta(n)$

**Parallelism:**  $T_1(n)/T_\infty(n) = \Theta(n)$



# Analysis of Nested Parallel Loops

```
// indices run from 0, not 1
cilk_for (int i=1; i<n; ++i) {
    cilk_for (int j=0; j<i; ++j) {
        double temp = A[i][j];
        A[i][j] = A[j][i];
        A[j][i] = temp;
    }
}
```

Span of outer loop  
control =  $\Theta(\lg n)$ .

Max span of inner loop  
control =  $\Theta(\lg n)$ .

Span of body =  $\Theta(1)$ .

**Work:**  $T_1(n) = \Theta(n^2)$

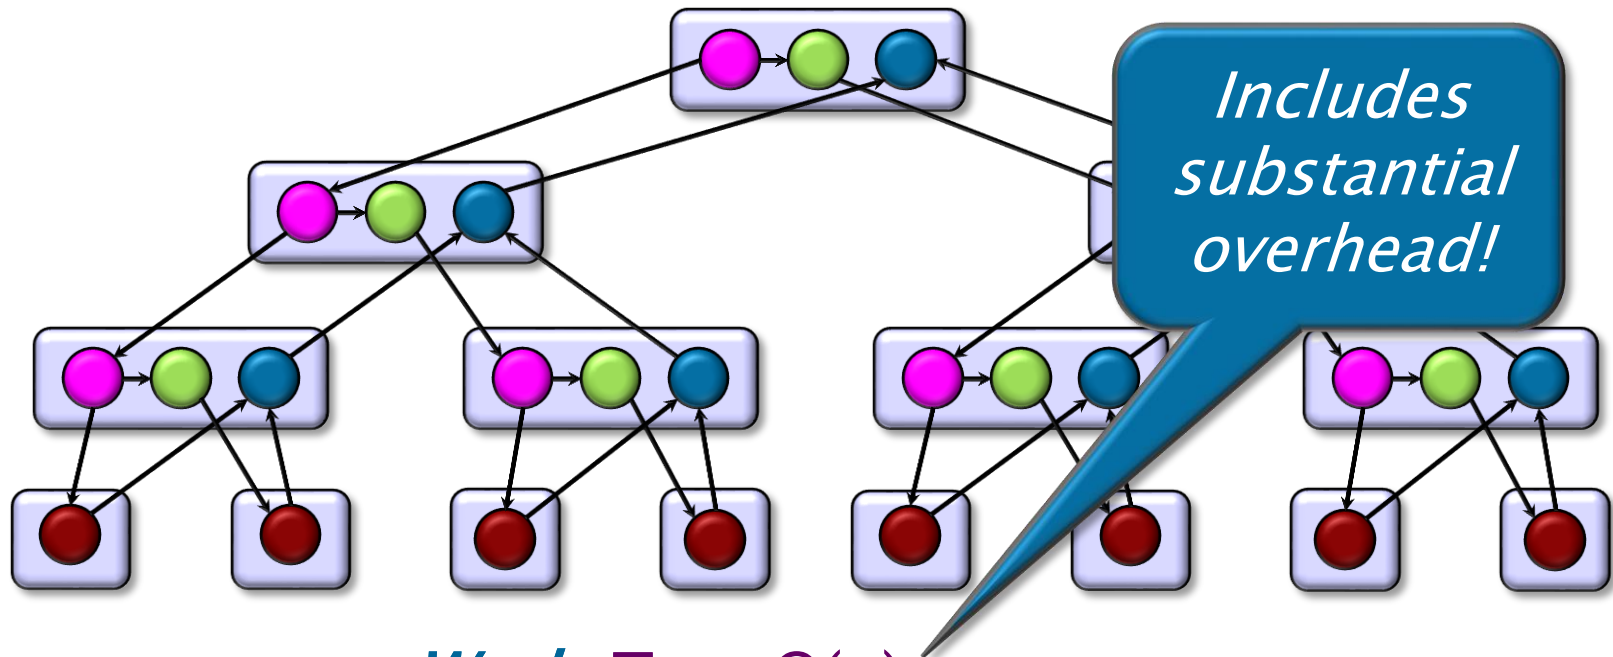
**Span:**  $T_\infty(n) = \Theta(\lg n)$

**Parallelism:**  $T_1(n)/T_\infty(n) = \Theta(n^2/\lg n)$

# A Closer Look at Parallel Loops

*Vector  
addition*

```
cilk_for (int i=0; i<n; ++i) {  
    A[i] += B[i];  
}
```



**Work:**  $T_1 = \Theta(n)$

**Span:**  $T_\infty = \Theta(\lg n)$

**Parallelism:**  $T_1 / T_\infty = \Theta(n / \lg n)$

# Coarsening Parallel Loops

```
#pragma cilk grainsize G
cilk_for (int i=0; i<n; ++i) {
    A[i] += B[i];
}
```

*Implementation  
with coarsening*

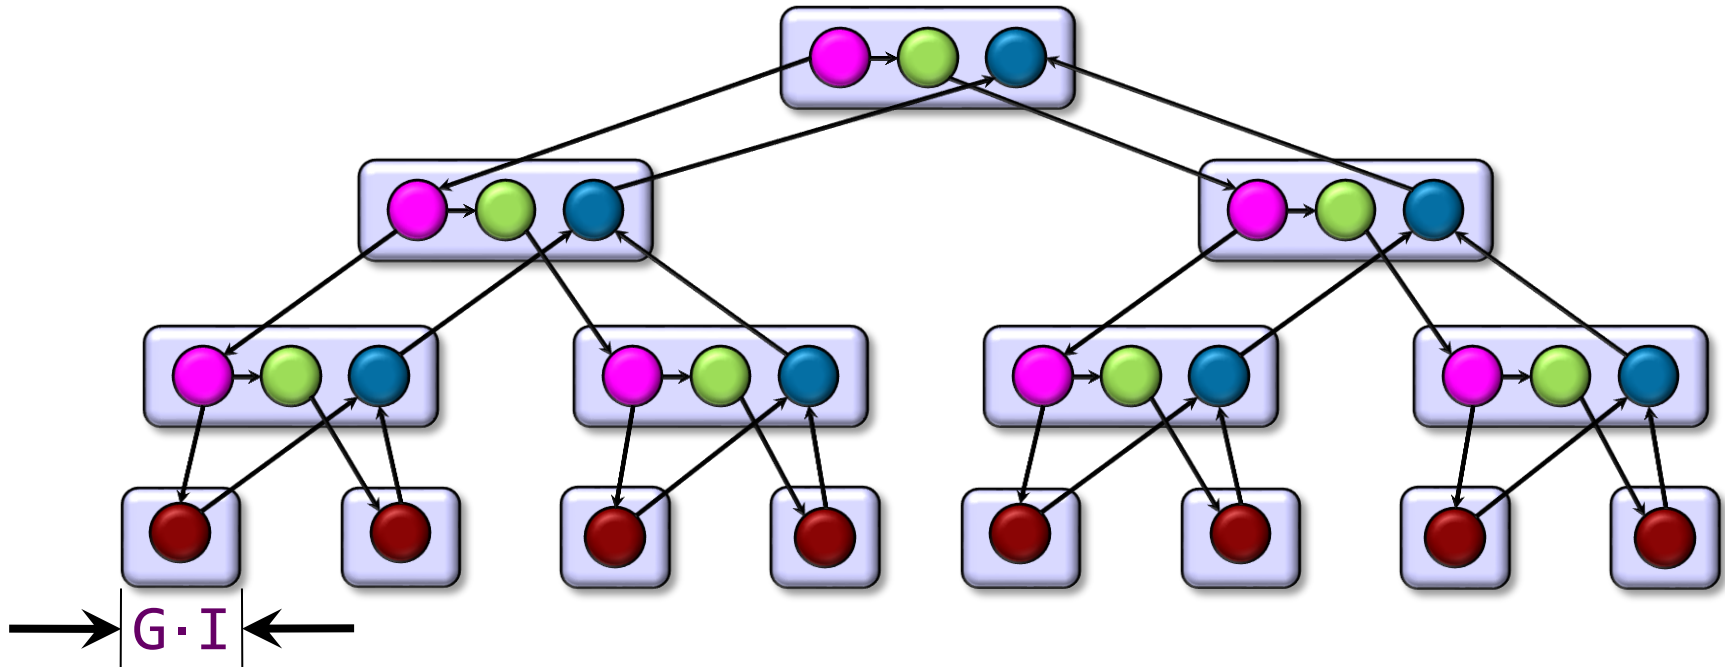
If a grainsize pragma is not specified, the Cilk runtime system makes its best guess to minimize overhead.

```
void recur(int lo, int hi) { //half open
    if (hi > lo + G) {
        int mid = lo + (hi - lo)/2;
        cilk_spawn recur(lo, mid);
                    recur(mid, hi);
        cilk_sync;
        return;
    }
    for (int i=lo; i<hi; ++i) {
        A[i] += B[i];
    }
}
...
recur(0, n);
```

# Loop Grain Size

*Vector  
addition*

```
#pragma cilk grainsize G
cilk_for (int i=0; i<n; ++i) {
    A[i] += B[i];
}
```

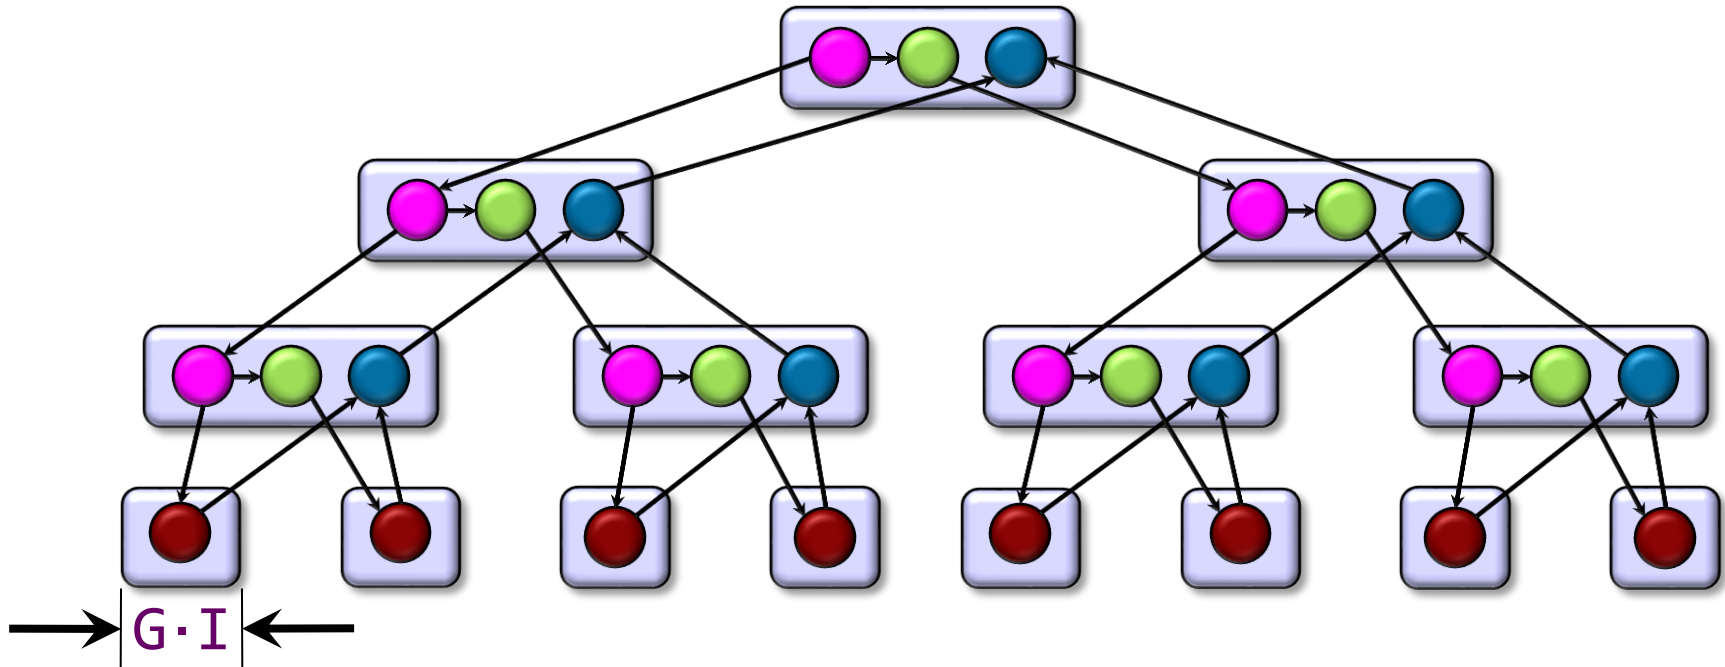


Let  $I$  be the time for one iteration of the loop body.  
Let  $S$  be the time to perform a spawn and return.

# Loop Grain Size

*Vector  
addition*

```
#pragma cilk grainsize G
cilk_for (int i=0; i<n; ++i) {
    A[i] += B[i];
}
```



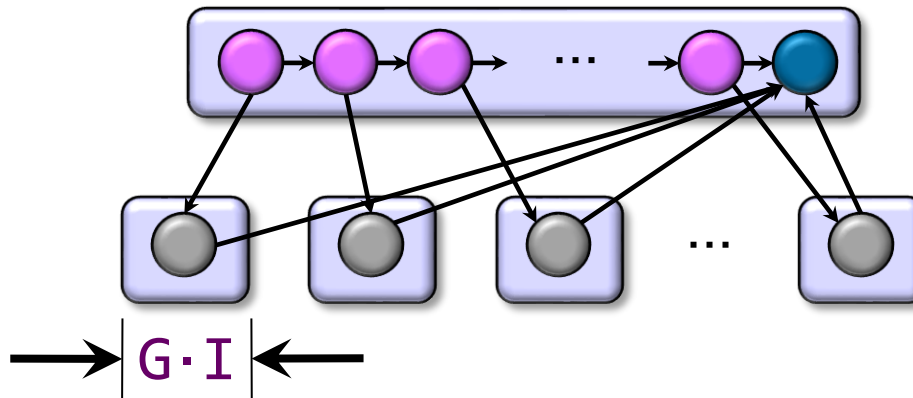
**Work:**  $T_1 = n \cdot I + (n/G - 1) \cdot S$

**Span:**  $T_\infty = G \cdot I + \lg(n/G) \cdot S$

Want  $G \gg S/I$   
and  $G$  small.

# Another Implementation

```
void vadd (double *A, double *B, int n) {  
    for (int i=0; i<n; i++) A[i] += B[i];  
}  
:  
for (int j=0; j<n; j+=G) {  
    cilk_spawn vadd(A+j, B+j, MIN(G,n-j));  
}  
cilk_sync;
```



Assume that  $G = 1$ .

*Work:*  $T_1 = \Theta(n)$

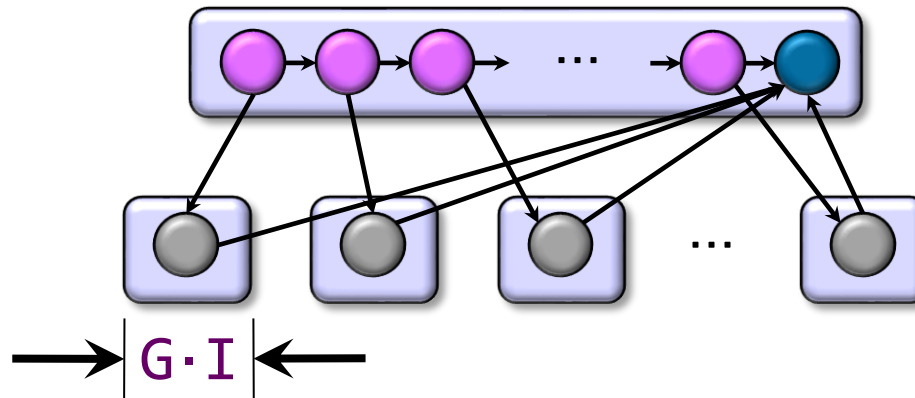
*Span:*  $T_\infty = \Theta(n)$

*Parallelism:*  $T_1/T_\infty = \Theta(1)$

PUNY!

# Another Implementation

```
void vadd (double *A, double *B, int n) {  
    for (int i=0; i<n; i++) A[i] += B[i];  
}  
:  
for (int j=0; j<n; j+=G) {  
    cilk_spawn vadd(A+j, B+j, min(G,n-j));  
}  
cilk_sync;
```



Choose  
 $G = \sqrt{n}$  to  
minimize.

Analyze in  
terms of  $G$ :

**Work:**  $T_1 = \Theta(n)$

**Span:**  $T_\infty \equiv \Theta(G + n/G) = \Theta(\sqrt{n})$

**Parallelism:**  $T_1/T_\infty = \Theta(\sqrt{n})$

# Quiz on Parallel Loops

**Question:** Let  $P$  be the number of workers on the system. How does the parallelism of the first code compare to that of the second? (Differences highlighted.)

```
#pragma cilk grainsize 1
cilk_for (int i=0; i<n; i+=32) {
    for (int j=i; j<min(i+32, n); ++j)
        A[j] += B[j];
}
```

```
#pragma cilk grainsize 1
cilk_for (int i=0; i<n; i+=n/P) {
    for (int j=i; j<min(i+n/P, n); ++j)
        A[j] += B[j];
}
```

**Work:**  $T_1 = \Theta(n)$

**Span:**  $T_\infty = \Theta(\lg(n/32) + 32)$   
 $= \Theta(\lg n)$

**Parallelism:**

$T_1 /$

**PUNY!**

Want  $T_1 / T_\infty \gg P$ .

**Work:**  $T_1 = \Theta(n)$

**Span:**  $T_\infty = \Theta(\lg P + P)$   
 $= \Theta(n/P)$

**Parallelism:**  $T_1 / T_\infty = \Theta(P)$



# Three Performance Tips

1. Minimize the **span** to maximize parallelism. Try to generate 10 times more parallelism than processors for near-perfect linear speedup.
2. If you have plenty of parallelism, try to trade some of it off to reduce **work overhead**.
3. Use **divide-and-conquer recursion** or **parallel loops** rather than spawning one small thing after another.

*Do this:*

```
cilk_for (int i=0; i<n; ++i) {  
    foo(i);  
}
```

*Not this:*

```
for (int i=0; i<n; ++i) {  
    cilk_spawn foo(i);  
}  
cilk_sync;
```

# And Three More

4. Ensure that `work/#spawns` is sufficiently large.
  - Coarsen by using function calls and **inlining** near the leaves of recursion, rather than spawning.
5. Parallelize **outer loops**, as opposed to inner loops, if you're forced to make a choice.
6. Watch out for *scheduling overheads*.

*Do this:*

```
cilk_for (int i=0; i<2; ++i) {  
    for (int j=0; j<n; ++j)  
        f(i,j);  
}
```

*Not this:*

```
for (int j=0; j<n; ++j) {  
    cilk_for (int i=0; i<2; ++i)  
        f(i,j);  
}
```